



Actor Model Concurrency in Erlang

Processes and their interaction

Walter Cazzola

Dipartimento di Informatica
Università degli Studi di Milano
e-mail: cazzola@di.unimi.it
twitter: @w_cazzola



Actor Model Concurrency

Traditional (Shared-State) Concurrency

Threads are the traditional way of offering concurrency

- the execution of the program is split up into concurrently running tasks;
- such tasks operate on shared memory

Several problems

- race conditions with update loss

T_1 (withdraw(10))	T_2 withdraw(10))	Balance
if (balance - amount >= 0)	if (balance - amount >= 0)	15 €
balance -= amount;	balance -= amount;	15 €
		5 €
		-5 €

- deadlocks

P_1	P_2
lock(A)	lock(B)
lock(B)	lock(A)

Erlang (and also Scala) takes a different approach to concurrency: the Actor Model.



Actor Model Concurrency

Overview

Each object is an actor:

- it has a mailbox and a behavior;
- actors communicate through messages buffered in a mailbox

Computation is data-driven, upon receiving a message an actor

- can send a number of messages to other actors;
- can create a number of actors; and
- can assume a different behavior for dealing with the next message in its mailbox

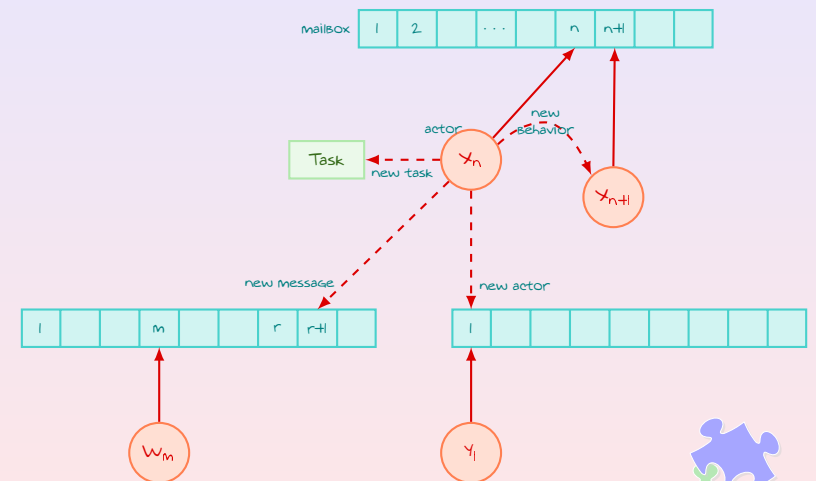
Note that,

- all communications are performed asynchronously;
 - the sender does not wait for a message to be received upon sending it;
 - no guarantees about the receiving order but they will eventually be delivered.
- there is no shared state between actors
 - information about internal state are requested/provided by messages;
 - also internal state manipulation happens through messages.
- actors run concurrently and are implemented as lightweight user-space threads



Actor Model Concurrency

Transaction Overview





Concurrency in Erlang Overview

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state

Erlang
concurrency

spawn
send
receive
scheduling
named actors

References

Three basic elements form the foundation for concurrency

- a built-in function (`spawn()`) to create new actors;
- an operator (`!`) to send a message to another actor; and
- a mechanism to pattern-match message from the actor's mailbox



Slide 5 of 15



Concurrency in Erlang Spawning New Processes.

Actor Model
Concurrency
in Erlang


Walter Cazzola

Concurrency
shared-state

Erlang
concurrency

spawn
send
receive
scheduling
named actors

References



Slide 6 of 15



Concurrency in Erlang My First Erlang Process.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state

Erlang
concurrency

spawn
send
receive
scheduling
named actors

References

```
-module(processes_demo).
-export([start/2, loop/2]).

start(N,A) -> spawn (processes_demo, loop, [N,A]).

loop(0,A) -> io:format("~p(~p) ~p~n", [A, self(), stops]);
loop(N,A) -> io:format("~p(~p) ~p~n", [A, self(), N]), loop(N-1,A).
```

```
1> processes_demo:start(7,a),processes_demo:start(5,b),processes_demo:start(3,c).
a(<0.73.0>) 7
b(<0.74.0>) 5
a(<0.73.0>) 6
c(<0.75.0>) 3
b(<0.74.0>) 4
<0.75.0>
a(<0.73.0>) 5
c(<0.75.0>) 2
b(<0.74.0>) 3
a(<0.73.0>) 4
c(<0.75.0>) 1
b(<0.74.0>) 2
a(<0.73.0>) 3
c(<0.75.0>) stops
b(<0.74.0>) 1
a(<0.73.0>) 2
b(<0.74.0>) stops
a(<0.73.0>) 1
a(<0.73.0>) stops
```

`self()` returns the PID of the process.



Slide 7 of 15



Concurrency in Erlang Sending a Message.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state

Erlang
concurrency

spawn
send
receive
scheduling
named actors

References

Every actor is characterized by:

- an address which identifies the actor and
- a **mailbox** where the delivered messages but not cleared yet are stored;

Messages are sorted on arrival time.

To send a message to an actor:

- has to know the address (pid) of the target actor;
- to send its address (pid) to the target with the message if a reply is necessary; and
- to use the `send (!)` primitive.

`Exp1 ! Exp2`

- `Exp1` must identify an actor;
- `Exp2` any valid Erlang expression; the result of the send expression is the one of `Exp2`;
- the sending never fails also when the target actor doesn't exist or is unreachable;
- the sending operation never block the sender.



Slide 8 of 15



Concurrency in Erlang

Receiving a Message.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state

Erlang
concurrency
spawn
send
receive
scheduling
named actors

References

The receiving operation uses pattern matching.

```
receive
  Any -> do_something(Any)
end
```


- the actor pick out of the mailbox the oldest message that matches Any.

```
receive
  {Pid, something} -> do_something(Pid)
end
```

- the actor tries to pick out the oldest message that matches {Pid, something};
- if it fails the actor is blocked waiting for such a message

```
receive
  Pattern1 [when GuardSeq1] -> Body1 ;
  ...
  Patternn [when GuardSeqn] -> Bodyn
[after Exprt -> Bodyt]
end
```

- rules definition and evaluation is quite similar to the functions;
- when no pattern matches the mailbox the actor waits instead of raising an exception;
- to avoid waiting forever the clause **after** can be used, after Expr_t ms the actor is woke up.



Slide 9 of 15



Concurrency in Erlang

Converting Some Temperatures.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state


Erlang
concurrency
spawn
send
receive
scheduling
named actors

References

```
-module(converter).
-export([t_converter/0]).

t_converter() ->
  receive
    {toF, C} -> io:format("~p °C is ~p °F~n", [C, 32+C*9/5]), t_converter();
    {toC, F} -> io:format("~p °F is ~p °C~n", [F, (F-32)*5/9]), t_converter();
    {stop} -> io:format("Stopping!~n");
    Other -> io:format("Unknown: ~p~n", [Other]), t_converter()
  end.
```

```
1> Pid = spawn(converter, t_converter, []).
<0.39.0>
2> Pid ! {toC, 32}.
32 °F is 0.0 °C
{toC,32}
3> Pid ! {toF, 100}.
100 °C is 212.0 °F
{toF,100}
4> Pid ! {stop}.
Stopping!
{stop}
5> Pid ! {toF, 100}. % once stopped a message to such a process is silently ignored
{toF,100}
```



Slide 10 of 15



Concurrency in Erlang

Calculating Some Areas.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state


Erlang
concurrency
spawn
send
receive
scheduling
named actors

References

```
-module(area_server).
-export([loop/0]).

loop() ->
  receive
    {rectangle, Width, Ht} ->
      io:format("Area of rectangle is ~p~n", [Width * Ht]),
      loop();
    {circle, R} ->
      io:format("Area of circle is ~p~n", [3.14159 * R * R]),
      loop();
    Other ->
      io:format("I don't know how to react to the message ~p~n", [Other]),
      loop()
  end.
```

```
1> Pid = spawn(fun area_server:loop/0).
<0.34.0>
2> Pid ! {rectangle, 30, 40}.
Area of rectangle is 1200
{rectangle,30,40}
4> Pid ! {circle, 40}.
Area of circle is 5026.544
{circle,40}
5> Pid ! {triangle, 22, 44}.
I don't know what the area of a {triangle,22,44} is
{triangle,22,44}
```



Slide 11 of 15



Concurrency in Erlang

Actor Scheduling in Erlang.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency
shared-state

Erlang
concurrency
spawn
send
receive
scheduling
named actors

References

Actors are not processes and are not dealt by the operating system

- the BEAM uses a preemptive scheduler;
- when an actor run for a too long period of time or when it enters a **receive** statement with no message available, the actor is halted and placed on a scheduling queue;

Actors and the rest of the system

- OS processes and actors have different schedulers and long running Erlang applications do not interfere with the execution of the OS processes (no one will become unresponsive)
- the BEAM supports symmetric multiprocessing (SMP)
 - i.e., it can run processes in parallel on multiple CPUs
 - but it cannot run lightweight processes (actors) in parallel on multiple CPUs.



Slide 12 of 15



Concurrency in Erlang

Timing the Spawning Process.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency

shared-state

Erlang

concurrency

spawn

send

receive

scheduling

named actors

References

```
-module(processes).
-export([max/1]).

max(N) ->
  Max = erlang:system_info(process_limit),
  io:format("Maximum allowed processes:~p~n",[Max]),
  statistics(runtime), statistics(wall_clock),
  L = for(1, N, fun() -> spawn(fun() -> wait() end) end),
  {_, Time1} = statistics(runtime), {_, Time2} = statistics(wall_clock),
  lists:foreach(fun(Pid) -> Pid ! die end, L),
  U1 = Time1 * 1000 / N, U2 = Time2 * 1000 / N,
  io:format("Process spawn time = ~p (~p) microseconds~n", [U1, U2]).

wait() -> receive die -> void end.

for(N, N, F) -> [F()];
for(I, N, F) -> [F()]|for(I+1, N, F)].
```

```
1> processes:max(20000).
Maximum allowed processes:32768
Process spawn time = 2.5 (3.4) microseconds
ok
2> processes:max(40000).
Maximum allowed processes:32768
=ERROR REPORT=== 8-Nov-2011::14:24:32 ===
Too many processes
...
[16:48]cazzola@surtur:~/lp/erlang-erl +P 100000
1> processes:max(50000).
Maximum allowed processes:100000
Process spawn time = 3.2 (3.74) microseconds
ok
```



Slide 13 of 15



Concurrency in Erlang

Giving a Name to the Actors.

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency

shared-state

Erlang

concurrency

spawn

send

receive

scheduling

named actors

References

Erlang provides a mechanism to render public the pid of a process to all the other processes.

- **register**(an_atom, Pid)
- **unregister**(an_atom)
- **whereis**(an_atom) -> Pid | **undefined**
- **registered**()

Once registered

- it is possible to send a message to it directly (name!msg).

```
-module(clock).
-export([start/2, stop/0]).

start(Time, Fun) -> register(clock, spawn(fun() -> tick(Time, Fun) end)).
stop() -> clock ! stop.

tick(Time, Fun) ->
  receive
  stop -> void
  after
  Time -> Fun(), tick(Time, Fun)
  end.
```

```
5> clock:start(5000, fun() -> io:format("TICK ~p~n",[erlang:now()]) end).
true
TICK 1320,769016,673190
TICK 1320,769021,678451
TICK 1320,769026,679120
7> clock:stop().
stop
```



Slide 14 of 15



References

Actor Model
Concurrency
in Erlang

Walter Cazzola

Concurrency

shared-state

Erlang

concurrency

spawn

send

receive

scheduling

named actors

References

- Gul Agha.
Actors: A Model of Concurrent Computation in Distributed Systems.
MIT Press, Cambridge, 1986.
- Joe Armstrong.
Programming Erlang: Software for a Concurrent World.
The Pragmatic Bookshelf, fifth edition, 2007.
- Francesco Cesarini and Simon J. Thompson.
Erlang Programming: A Concurrent Approach to Software Development.
O'Reilly, June 2009.



Slide 15 of 15